

## Lab 6 - Hive

---

- Submit your *own work* on time. No credit will be given if the lab is submitted after the due date.
  - Note that the completed lab should be submitted in .doc, .docx, .rtf or .pdf format only.
- 

This document is divided into two parts.

### 1. [Hive Practice Lab](#)

Just try to run through all the steps and see if they work properly for you.  
No need to submit this part.

We'll look at the following things:

- Create Database and tables
- Observe Warehouse and Metastore
- Run some more Hive Commands

### 2. [Hive Homework Lab](#)

You need to submit a document (not .zip) wherein I should be able to find all the instructions and commands.

Paste screenshots wherever applicable.

## Hive Practice Lab

In this practice lab, we'll do some analysis on the weather dataset.

1. Get the given simplified version of NCDC sample weather dataset and put the file in the local path `'/home/cloudera/cs523/input/weatherHive.txt'`
2. Now we'll load this local data into Hive's managed storage. Alternatively, try to put the input file in HDFS and see if you can load this HDFS data into Hive.
  - i. Just like in RDBMS, Hive organizes its data into tables. We create a table to hold the weather data using the CREATE TABLE statement:

```
CREATE TABLE records (year STRING, temperature INT, quality INT)
ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t';
```

The first line declares a records table with three columns: year, temperature, and quality. The type of each column must be specified, too. Hive expects there to be three fields in each row, corresponding to the table columns, with fields separated by tabs and rows by newlines.

After executing this create table statement, check the Metastore and the Warehouse directory and verify that "records" is created.

- ii. Next, populate Hive with the sample weather data.

```
LOAD DATA LOCAL INPATH '/home/cloudera/cs523/input/weatherHive.txt'
OVERWRITE INTO TABLE records;
```

Running this command tells Hive to put the specified local file in its warehouse directory. There is no attempt, for example, to parse the file and store it in an internal database format, because Hive does not mandate any particular file format.

Tables are stored as directories under Hive's warehouse directory, which is controlled by the `hive.metastore.warehouse.dir` property and defaults to `/user/hive/warehouse`.

Thus, the files for the `records` table are found in the `/user/hive/warehouse/records` directory on HDFS:

In this case, there is only one file, `weatherHive.txt` in that folder, but in general there can be more, and Hive will read all of them when querying the table.

- iii. Think : How will you change the above commands if you want to load the data from HDFS directory or if you want to create an external table?

3. Now that the data is in Hive, we can run a query against it:

```
hive> SELECT year, MAX(temperature) FROM records WHERE temperature !=  
9999 AND quality IN (0, 1, 4, 5) GROUP BY year;
```

This SQL query is unremarkable. It is a SELECT statement with a GROUP BY clause for grouping rows into years, which uses the MAX aggregate function to find the maximum temperature for each year group.

The remarkable thing is that Hive transforms this query into a job, which it executes on our behalf, then prints the results to the console. It is the ability to execute SQL queries against our raw data that gives Hive its power.

4. Now delete the above created table and check warehouse directory, what do you observe?

```
hive> DROP TABLE records;
```

5. Now repeat the above steps but this time create records as "EXTERNAL" table in hive. Note down what happens when you delete this external table.

6. No need to submit anything for this "practice lab" part.

---

### **Optional Part – Making Hive work with Avro**

Create your own example to understand how Avro works with Hive. Take help from this link:  
<https://cwiki.apache.org/confluence/display/Hive/AvroSerDe>

## Hive Homework Lab

In this HW lab, you need to analyze the city of Chicago employees' dataset.

- Create EXTERNAL (not managed) table in Hive for the given sample data set and find out some interesting facts from this data. For example, what's the maximum salary of employees in each department, etc.

**Show at least 3 different analysis with "limited" rows.**

### **What to submit :**

1. All the HQL statements used to complete the above requirements.
2. Screenshots of the step by step process also showing the output.

---

### **Important considerations for this HW**

In Hive, the most important part is how you create table for the HDFS data. So if you are not doing this step correctly then that means you missed the whole point of doing this lab! Remember, I'm not testing your SQL knowledge here!

You'll have to use a special SerDe for this dataset because you need to take care of the header row and the "comma" in name. You'll lose some points for not doing these things correctly. You are not allowed to modify the given .csv file!

If you plan to use a special SerDe then there might be an issue with datatypes. So always do a describe to check the datatypes of your created table.

Once your Hive table is properly created, I should be able to fire my own SQL query and it should give me correct results. You should never say that the way I've created this table, it'll only work with my 3 SQL queries!